

Modern Architectures & Optimization

Luca Tornatore - I.N.A.F. 

“Foundation of HPC” course




DATA SCIENCE &
SCIENTIFIC COMPUTING
2020-2021 @ Università di Trieste




Outline



First
things
first




Cache &
Memory



Branches



Pipelines




Loops



Outline



Avoid the
avoidable
inefficiencies



Loops
techniques




Prefetching

Introducing the example's framework

Let's suppose that


- 1) we have a distribution of random data points on a 2D plane which we subdivide in sub-regions using a grid.



Introducing the example's framework

Let's suppose that


- 1) we have a distribution of random data points on a 2D plane which we subdivide in sub-regions using a grid.
- 2) for each point p , we want to select all the grid cells whose center is closer to p than a given radius r , and to perform some operations accordingly to our search result.



Introducing the example's framework

Let's suppose that


- 1) we have a distribution of random data points on a 2D plane which we subdivide in sub-regions using a grid.
- 2) for each point p , we want to select all the grid cells whose center is closer to p than a given radius r , and to perform some operations accordingly to our search result.



Introducing the example's framework

We may consider to
use a nested loop
like this one →

Is there anything
you would change
in this loop?



```
for(p = 0; p < Np; p++)
```

```
    for(i = 0; i < Ng; i++)
```

```
        for(j = 0; j < Ng; j++)
```

```
            for(k = 0; k < Ng; k++)
```

```
            {
```

```
                dist = sqrt(
```

```
                    pow(x[p] - (double)i/Ng - half_size, 2) +
```

```
                    pow(y[p] - (double)j/Ng - half_size, 2) +
```

```
                    pow(z[p] - (double)k/Ng - half_size, 2));
```

```
                if(dist < R)
```

```
                    do something;
```

```
            }
```



| (1) Avoid expensive function calls



```
for(p = 0; p < Np; p++)
```

Some function calls are particularly expensive. Those include, among others, `sqrt()`, ...

Try to avoid them *if possible*.

```
    for(i = 0; i < Ng; i++)
        for(j = 0; j < Ng; j++)
            for(k = 0; k < Ng; k++)
                {
                    dist2 = pow(x[p] - (double)i/Ng - half_size, 2) +
                        pow(y[p] - (double)j/Ng - half_size, 2) +
                        pow(z[p] - (double)k/Ng - half_size, 2));

                    if(dist2 < R2)
                        do something;
                }
```




| (1) Avoid expensive function calls



```
for(p = 0; p < Np; p++)
```

Some function calls are particularly expensive. Those include, among others, `sqrt()`, `pow()`, ...

Try to avoid them *if possible*.

```
    for(i = 0; i < Ng; i++)
        for(j = 0; j < Ng; j++)
            for(k = 0; k < Ng; k++)
                {
                    dx = x[p] - (double)i/Ng - half_size;
                    dy = y[p] - (double)j/Ng - half_size;
                    dz = z[p] - (double)k/Ng - half_size;

                    dist2 = dx*dx + dy*dy + dz*dz;
                    if(dist2 < R2)
                        do something;
                }
```



| (1) Avoid expensive function calls



```
for(p = 0; p < Np; p++)
```

Some function calls are particularly expensive. Those include, among others, `sqrt()`, `pow()`, floating point division, ..
Try to avoid them if possible.

```
    for(i = 0; i < Ng; i++)
        for(j = 0; j < Ng; j++)
            for(k = 0; k < Ng; k++)
            {
                dx = x[p] - (double)i * Ng_inv - half_size;
                dy = y[p] - (double)j * Ng_inv - half_size;
                dz = z[p] - (double)k * Ng_inv - half_size;

                dist2 = dx*dx + dy*dy + dz*dz;
                if(dist2 < R2)
                    do something with sqrt(dist2);
            }
```



| (1) Avoid expensive function calls



```
(double)<i,j,k> * Ng_inv + half_size
```

was performed N^3+N^2+N times, always returning the same values.

Hoisting would save

$N(N^2+N^1+1)$ **mul**, **add** and **mem** accesses.

You can do better pre-computing the relevant values:

```
double ijk[Ng];  
for(i = 0; i < Ng; i++)  
    ijk[i] = i * Ng_inv + half_size
```

```
for(p = 0; p < Np; p++)
```

```
    for(i = 0; i < Ng; i++)  
        for(j = 0; j < Ng; j++)  
            for(k = 0; k < Ng; k++)  
            {  
                dx = x[p] - (double)i * Ng_inv - half_size;  
                dy = y[p] - (double)j * Ng_inv - half_size;  
                dz = z[p] - (double)k * Ng_inv - half_size;  
  
                dist2 = dx*dx + dy*dy + dz*dz;  
                if(dist2 < R2)  
                    do something with sqrt(dist2);  
            }
```



| (2) *Hoisting* of expressions

```
for(i = 0; i < Ng; i++) {  
    dx2 = x[p] - (double)i * Ng_inv - half_size;  
    dx2 = dx2*dx2;
```

(double)<i,j,k> * Ng_inv + half_size

was performed N^3+N^2+N times,
always returning the same values.
Hoisting would save
 $N(N^2+N^1+1)$ **mul**, **add** and **mem** accesses.

```
    for(j = 0; j < Ng; j++) {  
        dy2 = y[p] - (double)j * Ng_inv - half_size;  
        dy2 = dy2*dy2;  
        dist2_xy = dx2 + dy2;  
  
        for(k = 0; k < Ng; k++) {  
            dz = z[p] - (double)k * Ng_inv - half_size;  
            dist2 = dist2_xy + dz*dz;  
            if(dist2 < Rmax2)  
                do something with sqrt(dist2); } } }
```



| (2) *Hoisting* of expressions



You could do even better by pre-computing the relevant values:

```
double ijk[Ng];  
for(i = 0; i < Ng; i++)  
    ijk[i] = i * Ng_inv + half_size
```

```
for(i = 0; i < Ng; i++) {  
    dx2 = x[p] - Ng_inv[i] - half_size;  
    dx2 = dx2*dx2;  
  
    for(j = 0; j < Ng; j++) {  
        dy2 = y[p] - Ng_inv[j] - half_size;  
        dist2_xy = dx2 + dy2*dy2;  
  
        for(k = 0; k < Ng; k++) {  
            dz = z[p] - Ng_inv[k] - half_size;  
            dist2 = dist2_xy + dz*dz;  
            if(dist2 < Rmax2)  
                do something with sqrt(dist2); } } }
```



| (3) Clarify the variables' scope

All these variables are very local, there's no need for them to have a wider scope.

That will help you in writing the code, and *may* help the compiler in optimizing the stack and perhaps the registers usage.

```
for(int i = 0; i < Ng; i++) {  
    double dx2 = x[p] - (double)i * Ng_inv - half_size;  
    dx2 *= dx2;  
    for(j = 0; j < Ng; j++) {  
        double dy2 = y[p] - (double)j * Ng_inv - half_size;  
        double dist2_xy = dx2 + dy2*dy2;  
        for(k = 0; k < Ng; k++) {  
            double dz = z[p] - (double)k * Ng_inv - half_size;  
            double dist2 = dist2_xy + dz*dz;  
            if(dist2 < Rmax2)  
                do something with sqrt(dist2); } } }
```




| (4) Suggest what is important



```
double register Ng_inv = 1.0 / Ng;  
for(int i = 0; i < Ng; i++) {  
    double dx2 = x[p] - (double)i * Ng_inv - half_size;  
    dx2 *= dx2;
```

These variables are often calculated and reused subsequently.

```
    for(j = 0; j < Ng; j++) {  
        double dy2 = y[p] - (double)j * Ng_inv - half_size;  
        dy2 *= dy2;  
        double register dist2_xy = dx2 + dy2;
```

Keeping a register dedicated to them may be useful.


```
    for(k = 0; k < Ng; k++) {  
        double register dz = z[p] - (double)k * Ng_inv - ...;  
        double register dist2 = dist2_xy + dz*dz;
```

Note: this is a suggestion, the compiler, after analyzing the code, may decide differently


```
    if(dist2 < Rmax2)  
        do something with sqrt(dist2); } } }
```



Outline




Avoid the
avoidable
inefficiencies



Loops
techniques



Prefetching



Optimizing cache access in loops

Loop classification

$$A_I = \frac{f(n)}{n}$$

Arithmetic Intensity: the ratio between the number of performed operations and the amount of the data.

1. $O(N) / O(N)$
optimization potential limited
2. $O(N^2) / O(N^2)$
some more opportunities for opt.
3. $O(N^3) / O(N^2)$
significant optimization potential



Cache access in loops: $O(N)/O(N)$

Example

1-level loops: Scalar products, vector additions, sparse matrix-vector multiplication

Inevitably memory-bound for very large N ; in general, improvements come from *avoiding unnecessary operations* and/or *repeated memory accesses*, and increasing **data reuse**

[check the room for loops fusion]

$O(N) / O(N)$


```
for(int j=0; j<2; j++)  
    A[i] = B[i] × C[i]
```



```
for(int j=0; j<2; j++)  
    Q[i] = B[i] + D[i]
```

```
for(int j=0; j<2; j++)  
{  
    A[i] = B[i] × C[i]  
    Q[i] = B[i] + D[i]  
}
```

Loop fusion: in the version on the right, B is recalled from memory only once.




Cache access in loops: $O(N^2)/O(N^2)$

Example


2-levels loops: dense matrix-vector mul, matrix transpos., matrix add, ...

Improvements comes again from increasing *data reuse*, exploiting *locality* and *avoiding unnecessary* operations and memory accesses.

```
for(int i=0; i < N; i++)  
    for(int j=0; j<N; j++)  
        C[i] += A[i][j] * B[j];
```



→ $3 \times N^2$ memory accesses



Cache access in loops: $O(N^2)/O(N^2)$

Step 1:

Avoid unnecessary loads /stores


```
for(int i=0; i < N; i++)  
    for(int j=0; j<N; j++)  
        C[i] += A[i][j] * B[j];
```

→

```
for(int i=0; i < N; i++) {  
    c_temp = C[i];  
    for(int j=0; j < N; j++)  
        c_temp += A[i][j] * B[j];  
    C[i] = c_temp; }
```

Now it is clearer for the compiler that **C[i]** need to be loaded and stored only 1 time

→ $2 \times N^2 + N$ memory accesses



Cache access in loops: $O(N^2)/O(N^2)$

Step 2:

Unroll outern loop and *fuse* in the inner loop; there is potential for *vectorisation*.

```
for(int i=0; i < N; i++)  
    for(int j=0; j<N; j++)  
        C[i] += A[i][j] * B[j];
```

→

```
for(int i=0; i < N; i += m)      N×N/m  
    for(j = 0; j < N; j++){  
        b_temp = B[j];  
        C[i]   += A[i][j] * b_temp;  
        C[i+1] += A[i+1][j] * b_temp;  
        ...  
        C[i+m] += A[i+m][j] * b_temp; }  
N
```

Diagram illustrating the transformation of the loop structure. The original nested loop is transformed into an unrolled version where the outer loop increments by m . The inner loop is fused, and the vector C is updated in a vectorized manner. The transformation is annotated with $N \times N$ and N indicating the dimensions of the data structures involved.

→ $N^2 \times (1 + 1/m) + N$



| Note: unrolling and register spill

Using a too large m in the previous example while the target CPU does not have enough registers to keep all the needed operands results in a “code bloating”.

In this case, the CPU has to spill registers' content to cache and viceversa, slowing down the computation.

→ learn to inspect the compiler's *log*

A too much involute and obscure loop body may hamper the compiler to effectively perform *unroll* & *jam* optimizations targeted to the CPU it runs on.

→ hand code effort to clarify the code

→ hints / directives to the compiler

(directives are generally not portable across different compilers)



| Cache access in loops: $O(N^2)/O(N^2)$


Sometimes no magic wand can cure the fact that you have to access N^2 memory locations.

For instance: in matrix transpose you have to access all the source matrix and all the destination matrix once.




Unroll & Jam strategy can bring benefits as long as the cache can hold N lines.

An L_C -way unrolling is too much aggressive and may easily result in register pressure.


Loop blocking is a good strategy that does not save memory loads but increase dramatically the cache hit ratio



Cache access in loops: $O(N^2)/O(N^2)$



Step 3:
Fully exploit locality
of referenced data;
cut TLB misses by
accessing 2D arrays
by blocks



| Loop unrolling

Loop unrolling is a fundamental code transformation which usually helps significantly in improving your code performance:

- It reduces the loop overhead (counter update, branching)
- It exposes *critical data path* and dependencies
- It helps in exploiting ILP, especially in case of memory aliasing



| Cache access in loops: $O(N^3)/O(N^2)$


These algorithms (ex: matrix-matrix multiplication or dense matrix diagonalization) are very good candidates for optimizations that lead flop/s performance very close to the theoretical peak (in fact, MMM is at the core of **linpack**).

Blocking, unroll&jam + vectorization of operations, reorganization of ops to exploit CPU's pipelines and out-of-order capability, are all used by extremely specialized libraries.

→ It is a brilliant idea to link those library instead of developing your own algorithm, unless some very special needs must be met.



Outline



Avoid the
avoidable
inefficiencies



Loops
techniques



Prefetching



At the right moment, at the right place

We know that waiting for data and instructions is a major performance killer.

Modern CPUs have the capability of pre-emptively bring from memory into cache levels data that **will be needed shortly afterwards**.

They can do that following some speculative algorithm based on the current execution flow and assuming spatial locality and temporal locality.

Both *data* and *instructions* can be pre-fetched.

Pre-fetching may be both hardware-based and software-based (typically the compiler insert pre-fetching instructions at compile-time).



At the right moment, at the right place

From the point of view of the programmer, there are 2 possible ways to deal with prefetching:

EXPLICIT

you explicitly insert a pre-fetching directive.

Very difficult to be achieved effectively: the directive must be inserted timely but not too early (data eviction) or too late (load latency).

INDUCED

you consciously arrange data layout and execution flow so that to make it obvious to the compiler what to prefetch.



Explicit prefetching

This is a standard binary search implementation.

Find the median element

Define the next search

```
int mybsearch(int *data, int N, int Key)
{
    int register low = 0;
    int register high = N;
    int register mid;

    while(low <= high) {
        mid = (low + high) / 2;

        if(data[mid] < Key)
            low = mid + 1;
        else if(data[mid] > Key)
            high = mid-1;
        else
            return mid;
    }
    return -1;
}
```



Explicit prefetching

We can make it better by simply making sure that the element to be compared for (the mid) is in the cache when requested

```
int mybsearch(int *data, int N, int Key)
{
    int register low = 0;
    int register high = N;
    int register mid;

    while(low <= high) {
        mid = (low + high) / 2;

        if(data[mid] < Key)
            low = mid + 1;
        else if(data[mid] > Key)
            high = mid-1;
        else
            return mid;
    }
    return -1;
}
```



Explicit prefetching

We can make it better by simply making sure that the element to be compared for (the `mid`) is in the cache when requested

```
int mybsearch(int *data, int N, int Key)
{
    int register low = 0;
    int register high = N;
    int register mid;


    while(low <= high) {
        mid = (low + high) / 2;
        __builtin_prefetch (&data[(mid + 1 + high)/2], 0, 3);
        __builtin_prefetch (&data[(low + mid - 1)/2], 0, 3);

        if(data[mid] < Key)
            low = mid + 1;
        else if(data[mid] > Key)
            high = mid-1;
        else
            return mid;
    }
    return -1; }
```




Explicit prefetching

```
luca@GGG:~/code/HPC_LECTURES/prefetching% ./prefetching off  
performing 13421772 lookups on 134217728 data..  
set-up data.. set-up lookups..  
start cycle.. time elapsed: 20.7534  
luca@GGG:~/code/HPC_LECTURES/prefetching% ./prefetching on  
performing 13421772 lookups on 134217728 data with prefetching enabled..  
set-up data.. set-up lookups..  
start cycle.. time elapsed: 12.6204
```



Prefetching

Explicit prefetching

Samples: 71K of event 'cpu/mem-loads,ldlat=30/P', Event count (approx.): 13901140

Overhead	Samples	Memory access
----------	---------	---------------

71,08%	42196	Local RAM hit
--------	-------	---------------

24,14%	17022	LFB hit
--------	-------	---------

4,11%	10967	L3 hit
-------	-------	--------

0,63%	1714	L1 hit
-------	------	--------

0,02%	75	L2 hit
-------	----	--------

0,01%	15	L3 miss
-------	----	---------

0,00%	1	Uncached hit
-------	---	--------------

Samples: 61K of event 'cpu/mem-loads,ldlat=30/P', Event count (approx.): 11720387

Overhead	Samples	Memory access
----------	---------	---------------

68,74%	29450	LFB hit
--------	-------	---------

27,04%	28208	L1 hit
--------	-------	--------

2,72%	909	Local RAM hit
-------	-----	---------------

1,29%	2983	L3 hit
-------	------	--------

0,20%	346	L2 hit
-------	-----	--------



| Explicit prefetching

Usage of direct prefetching directive is highly uncertain, since it is difficult to spot the exact point – both in the code and in the execution – where to place them (also because your C code is different than the generated assembly code).

Moreover, the “exact point” is very likely dependent on the system you run on, and then it is susceptible to change significantly.

It is normally much safer to re-organize your code so to have **prefetching by pre-loading**.



Prefetching by moral suasion

Let's discuss together
this very simple example
before putting the hands
on the code you find in
the `git`

```
elem a = elements[0]
for ( i = 0; i < 4*N_4; i+= 4 )
{
    elem e = elem[i+4]; // non-blocking miss
    elem b = elem[i+1]; // possible cache-hit
    elem c = elem[i+2]; // possible cache-hit
    elem d = elem[i+3]; // possible cache-hit
    Elaborate(a);
    Elaborate(b);
    Elaborate(c);
    Elaborate(d);
    a = e;
}
```



You find code snippets with different flavours of prefetching-by-preloading technique on our GitHub, with some comments about compilation.

Compile and run them with different options (and possibly different compilers) and try to understand what happens on your laptop and/or on HPC facility.

```
for ( i = 0; i < N; i++ )  
    sum += array[ i ];
```

that's all, have fun



“So long
and thanks
for all the fish”